

Ljubljana, March 17th, 2021

Dear Editor,

We would like to submit our manuscript on a modular Python toolbox for t-SNE dimensionality reduction and embedding.

Science, according to Wikipedia, is “a systematic enterprise that builds and organizes knowledge”. One of the earliest means of information organization is a map. Today, an important subfield of data science strives to map multidimensional data onto a two-dimensional plane to expose the structure, relations, provide predictions of functions, and uncover temporal trends. A prominent example of such a technique is t-SNE. For instance, t-SNE can embed expression-profiled single cells into a two-dimensional map, exposing regions of same-typed cells or charting cell development. Over 1,400 articles from Nature journals alone recently used t-SNE for data visualization. Nature Methods has published a number of papers that use t-SNE, including those with a t-SNE depiction of single-cell data maps. t-SNE visualizations have appeared on the magazines’ covers and have been embraced by the entire data science community.

In the paper, we report on openTSNE, a Python-based open-source library for t-SNE visualization. openTSNE runs orders of magnitudes faster than comparable Python-based implementations (e.g., scikit-learn) and can handle data sets containing millions of data points. t-SNE has been recently criticized for poor scalability and susceptibility to batch effects. Our proposed implementation addresses these concerns and includes the lately-proposed methodological advancements, which we review in the paper’s online methods section.

Additionally, openTSNE is currently the only t-SNE library that can place new data points into a constructed embedding. The paper shows that such embedding of new data effectively mitigates batch effects and is an approach for visual prediction.

Our t-SNE implementation is open, available on GitHub, and fosters extensibility and experimentation. The library has already attracted significant interest within the Python community; it gained 758 GitHub stars, on par with prominent data science packages in Python, such as scanpy for single-cell analysis (854 stars).

We are submitting an original manuscript that has not been considered for publication before. Lin Tang would be an excellent choice for Editor for this manuscript because of his broad knowledge in the field and work in single-cell analytics. For reviewers, we kindly propose Dmitry Kobak (U Tuebingen), an author of recently proposed t-SNE tricks that we implemented in the library, Barbara di Camillo (U Padova), single-cell data scientists, or Cagatay Turkay (U Warwick), a specialist in explainable data visualizations.

Yours faithfully,

Pavlin G. Poličar, Martin Stražar, Blaž Zupan